✓ **Congratulations! You passed!**

| **Grade** | **Latest Submission** | **To pass** 80% or | **Go to next item** |
| received 100% | **Grade** 100% | higher | |

---

1.   `1 / 1 point`

The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- ⦿ The state contains numbers such as position and velocity that are continuous valued.
- ○ The state has multiple numbers rather than only a single number (such as position in the x-direction)
- ○ The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ○ The reward contains numbers that are continuous valued

✓ **Correct**
   That's right!

2.   `1 / 1 point`

In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is y = _____?

- ○ $y = R(s')$ where s' is the state you get to after taking action a in state s
- ○ $y = R(s)$
- ○ $y = \max\limits_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ⦿ $y = R(s) + \gamma \max\limits_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

✓ **Correct**

3.   `1 / 1 point`

You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

- ☑ What an accomplishment -- you made it!

✓ **Correct**

- ☑ You deserve to celebrate!

✓ **Correct**

- ☑ The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

✓ **Correct**

- ☑ Andrew sends his heartfelt congratulations to you!

✓ **Correct**